
INTRO TO DATA SCIENCE

FINAL PROJECT KICKOFF

Rob Hall

DAT9 SF // August 21, 2014

FINAL PROJECT KICKOFF

Final Projects in Brief:

- Select a dataset
- Formulate a hypothesis / problem statement
 - What question do you want to answer? / What do you want to predict?
 - Why is it useful to do so?
- Wrangle the data
- Build at least one model
- Test that model
- Write up your project and what you learned
- Present your project to the class (last week of class)

FINAL PROJECT KICKOFF

What is the purpose of doing a final project?

- Opportunity to apply everything you will have learned in this course
- Opportunity for your curiosity to come out and play
- Required to receive a “Letter of Completion” for this course
- Becomes part of your portfolio when seeking employment in data science
- There will be beer

FINAL PROJECT KICKOFF

Sources of inspiration:

- Your work
- Cool projects you see online
- Stanford's CS 229 projects:
<http://cs229.stanford.edu/projects2013.html>
<http://cs229.stanford.edu/projects2012.html>
- Each other

FINAL PROJECT KICKOFF

The “Elevator Pitch” (due Tuesday 9/2 in class)

- A one-paragraph write up of the project you propose
- A concise (<90 seconds) verbal pitch of your project to the rest of the class
- Include:
 - A concise statement of the goal of your project
 - Where you will get the data / what dataset you plan to use
 - What type of machine learning problem this is (from our 2x2)
 - Why you think this is a cool project
- *NOTE: This does not lock you into that project. You can change your idea and your ML technique as we learn more. The point is to get moving on it :-)*