INTRO TO DATA SCIENCE
LECTURE 8: LOGISTIC REGRESSION

Rob Hall
DAT9 SF // September 4, 2014

LAST TIME:

– INTRO TO BAYESIAN PROBABILITY
– NAIVE BAYESIAN CLASSIFICATION

QUESTIONS?

LECTURE:
I. LOGISTIC REGRESSION
II. OUTCOME VARIABLES
III. ERROR TERMS
IV. INTERPRETING RESULTS

LAB: IMPLEMENTING LOGISTIC REGRESSION IN PYTHON

IN-CLASS EXERCISE: FINAL PROJECT ELEVATOR PITCHES

I. LOGISTIC REGRESSION

	continuous	*categorical*
supervised	???	???
unsupervised	???	???

	continuous	**categorical**
supervised	regression	classification
unsupervised	dimension reduction	clustering

Q: What is **logistic regression***?*

Q: What is **logistic regression***?*

A: A generalization of the linear regression model to classification problems.

In linear regression, we used a set of covariates (independent variables) to predict the value of a <u>continuous</u> outcome variable.

In linear regression, we used a set of covariates (independent variables) to predict the value of a <u>continuous</u> outcome variable.

In logistic regression, we use a set of covariates to predict probabilities of (binary) class membership.

In linear regression, we used a set of covariates (independent variables) to predict the value of a <u>continuous</u> outcome variable.

In logistic regression, we use a set of covariates to predict <u>probabilities</u> of (binary) class membership.

These probabilities are then <u>mapped to class labels</u>, thus solving the classification problem (categorical output).

A motivating problem:

The following figure shows 30 day mortality in a sample of septic patients as a function of their baseline APACHE II score. Patients are coded as 1 or 0 depending on whether they are dead or alive in 30 days, respectively.

How can we predict death from baseline APACHE II score in these patients?

Q: How can we predict death from baseline APACHE II score in these patients?

Let p(x) be the probability that a patient with score x will die within 30 days.

Well, linear regression would not work well here, because it could produce probabilities less than zero or greater than one. Also, one new value could great change our model...

*So, what can we do instead of linear regression?*

30 Day Mortality in Patients with Sepsis

LOGISTIC REGRESSION

probability of belonging to class

value of independent variable

NOTE

Probability predictions look like this.

This function fits our problem much better:

$$0 \le h_\theta(x) \le 1$$

In other words, our classifier will output values between 0 and 1. It asymptotically approaches 0 and 1.

This is called the Sigmoid Function, or the _Logistic_ Function (synonymous)

*This function fits our problem much better:*

$$0 \leq h_\theta(x) \leq 1$$

In other words, our classifier will output values between 0 and 1. It asymptotically approaches 0 and 1.

This is called the Sigmoid Function, or the <u>*Logistic*</u> *Function (synonymous)*

NOTE

This function gives Logistic Regression its name!

The logistic function:

$$F(t) = \frac{e^t}{e^t + 1} = \frac{1}{1 + e^{-t}}$$

Notice that f(t) = 0.5 when t = 0

f(t) >= 0.5 when t > 0

f(t) <= 0.5 when t < 0

Suppose we predict class 1 when f(t) >= 0.5 and class 0 when f(t) < 0.5

$$F(t) = \frac{e^t}{e^t + 1} = \frac{1}{1 + e^{-t}}$$

So, if the t in the logistic function is a linear function of an explanatory variable x, or a linear combination of explanatory variables, the logistic function becomes:

$$F(x) = \frac{1}{1 + e^{-(\beta_0 + \beta_1 x)}}$$

$$F(t) = \frac{e^t}{e^t + 1} = \frac{1}{1 + e^{-t}}$$

So, if the t in the logistic function is a linear function of an explanatory variable x, or a linear combination of explanatory variables, the logistic function becomes:

$$F(x) = \frac{1}{1 + e^{-(\beta_0 + \beta_1 x)}}$$

Does that exponent look familiar...?

$$F(x) = \frac{1}{1 + e^{-(\beta_0 + \beta_1 x)}}$$

B_0	B_1
- 4	0.4
- 8	0.4
- 12	0.6
- 20	1.0

When $B_0 + B_1 x = 0$, then $F(x) = 0.5$, which is the inflection point on all these curves.

Going back to our example of patient survival given a sepsis test score: Data that has a sharp cut off point between the two classes (living / dying) should have a large value of B_1.

Going back to our example of patient survival given a sepsis test score: Data that has a lengthy transition between the two classes (living / dying) should have a small value of B_1.

PROBABILITIES

probability of belonging to class

value of independent variable

NOTE

Probability predictions look like this.

CLASS LABELS

class label

value of independent variable

NOTE

Probabilities are "snapped" to class labels (e.g. by thresholding at 50%).

The logistic regression model is an extension of the linear regression model, with a couple of important differences.

The logistic regression model is an extension of the linear regression model, with a couple of important differences.

The first difference is in the outcome variable.

The logistic regression model is an extension of the linear regression model, with a couple of important differences.

The first difference is in the outcome variable.

The second difference is in the error term.

II. OUTCOME VARIABLES

The key variable in any regression problem is the **conditional mean** *of the outcome variable y given the value of the covariate x:*

$$E(y|x)$$

*The key variable in any regression problem is the **conditional mean** of the outcome variable y given the value of the covariate x:*

$$E(y|x)$$

In linear regression, we assume that this conditional mean is a linear function taking values in (-∞, +∞):

$$E(y|x) = \alpha + \beta x$$

In logistic regression, we've seen that the conditional mean of the outcome variable takes values only in the unit interval [0, 1].

In logistic regression, we've seen that the conditional mean of the outcome variable takes values only in the unit interval [0, 1].

The first step in extending the linear regression model to logistic regression is to map the outcome variable $E(y|x)$ into the unit interval.

In logistic regression, we've seen that the conditional mean of the outcome variable takes values only in the unit interval [0, 1].

The first step in extending the linear regression model to logistic regression is to map the outcome variable $E(y|x)$ into the unit interval.

Q: How do we do this?

A: By using a transformation called the **logistic function***:*

$$E(y|x) = \pi(x) = \frac{e^{\alpha+\beta x}}{1+e^{\alpha+\beta x}}$$

A: By using a transformation called the **logistic function***:*

$$E(y|x) = \pi(x) = \frac{e^{\alpha+\beta x}}{1+e^{\alpha+\beta x}}$$

We've already seen what this looks like:

*A: By using a transformation called the* **logistic function**:

$$E(y|x) = \pi(x) = \frac{e^{\alpha+\beta x}}{1+e^{\alpha+\beta x}}$$

We've already seen what this looks like:

NOTE

For any value of x, y is in the interval [0, 1]

This is a nonlinear transformation!

The **logit function** *is an important transformation of the logistic function. Notice that it returns the linear model!*

$$g(x) = ln\left(\frac{\pi(x)}{1-\pi(x)}\right) = \alpha + \beta x$$

*The **logit function** is an important transformation of the logistic function. Notice that it returns the linear model!*

$$g(x) = ln(\frac{\pi(x)}{1-\pi(x)}) = \alpha + \beta x$$

*The logit function is also called the **log-odds** function.*

The **logit function** *is an important transformation of the logistic function. Notice that it returns the linear model!*

$$g(x) = ln\left(\frac{\pi(x)}{1-\pi(x)}\right) = \alpha + \beta x$$

NOTE

This name hints at its usefulness in interpreting our results.

We will see why shortly.

The logit function is also called the **log-odds function**.

III. ERROR TERMS

The second difference between linear regression and the logistic regression model is in the error term.

The second difference between linear regression and the logistic regression model is in the error term.

One of the key assumptions of linear regression is that the error terms follow independent Gaussian distributions with zero mean and constant variance:

$$\epsilon \sim N(0, \sigma^2)$$

In logistic regression, the outcome variable can take only two values: 0 or 1.

In logistic regression, the outcome variable can take only two values: 0 or 1.

It's easy to show from this that instead of following a Gaussian distribution, the error term in logistic regression follows a Bernoulli distribution:

$$\epsilon \sim B(0, \pi(1 - \pi))$$

In logistic regression, the outcome variable can take only two values: 0 or 1.

It's easy to show from this that instead of following a Gaussian distribution, the error term in logistic regression follows a Bernoulli distribution:

$$\epsilon \sim B(0, \pi(1 - \pi))$$

NOTE

This is the same distribution followed by a coin toss.

Think about why this makes sense!

These two key differences define the logistic regression model, and they also lead us to a kind of unification of regression techniques called **generalized linear models**.

These two key differences define the logistic regression model, and they also lead us to a kind of unification of regression techniques called **generalized linear models***.*

Briefly, GLMs generalize the distribution of the error term, and allow the conditional mean of the response variable to be related to the linear model by a **link function***.*

In the present case, the error term follows a Bernoulli distribution, and the logit is the link function that connects us to the linear predictor.

In the present case, the error term follows a Bernoulli distribution, and the logit is the link function that connects us to the linear predictor.

$$g(x) = ln\left(\frac{\pi(x)}{1-\pi(x)}\right) = \alpha + \beta x$$

In the present case, the error term follows a Bernoulli distribution, and the logit is the link function that connects us to the linear predictor.

$$g(x) = ln(\frac{\pi(x)}{1-\pi(x)}) = \alpha + \beta x$$

Since the Bernoulli distribution and the logit function share a common parameter π, we say that the logit is the **canonical link function** *for the Bernoulli distribution.*

In the present case, the error term follows a Bernoulli distribution, and the logit is the link function that connects us to the linear

NOTE

This terminology is just FYI!

$$g(x) = ln(\frac{\pi(x)}{1-\pi(x)}) = \alpha + \beta x$$

Since the Bernoulli distribution and the logit function share a common parameter π, we say that the logit is the **canonical link function** *for the Bernoulli distribution.*

IV. INTERPRETING RESULTS

In linear regression, the parameter β represents the change in the response variable for a unit change in the covariate.

In linear regression, the parameter β represents the change in the response variable for a unit change in the covariate.

In logistic regression, β represents the change in the logit function for a unit change in the covariate.

In linear regression, the parameter β represents the change in the response variable for a unit change in the covariate.

In logistic regression, β represents the change in the logit function for a unit change in the covariate.

Interpreting this change in the logit function requires another definition first.

*The **odds** of an event are given by the ratio of the probability of the event by its complement:*

$$O(x = 1) = \frac{\pi(1)}{(1 - \pi(1))}$$

*The **odds** of an event are given by the ratio of the probability of the event by its complement:*

$$O(x = 1) = \frac{\pi(1)}{(1 - \pi(1))}$$

*The **odds ratio** of a binary event is given by the odds of the event divided by the odds of its complement:*

$$OR = \frac{O(x=1)}{O(x=0)} = \frac{\pi(1)/[1 - \pi(1)]}{\pi(0)/[1 - \pi(0)]}$$

Substituting the definition of π(x) into this equation yields (after some algebra),

$$OR = e^{\beta}$$

Substituting the definition of π(x) into this equation yields (after some algebra),

$$OR = e^{\beta}$$

This simple relationship between the odds ratio and the parameter β is what makes logistic regression such a powerful tool.

Q: So how do we interpret this?

Q: So how do we interpret this?

A: The odds ratio of a binary event gives the increase in likelihood of an outcome if the event occurs.

Suppose we are interested in mobile purchase behavior. Let y be a class label denoting purchase/no purchase, and let x denote a mobile OS (for example, iOS).

Suppose we are interested in mobile purchase behavior. Let y be a class label denoting purchase/no purchase, and let x denote a mobile OS (for example, iOS).

In this case, an odds ratio of 2 (eg, $\beta = log(2)$) indicates that a purchase is twice as likely for an iOS user as for a non-iOS user.

LAB: LOGISTIC REGRESSION